

Docker + MongoDB + Flask Setup Guide

This guide explains how to prepare, install, and test a Flask application connected to MongoDB using Docker.

Preparation

1. Download the Source Code

Download the project source code from the provided Google Drive folder.

2. Install Docker (macOS)

Follow Docker Desktop installation for Mac.

Alternatively, using command line:

```
sudo hdiutil attach Docker.dmg
sudo /Volumes/Docker/Docker.app/Contents/MacOS/install
sudo hdiutil detach /Volumes/Docker
```

Check Docker version:

```
docker --version
```

3. Install Kubernetes Tools (Optional)

Install **kubect**l and **minikube** using Homebrew:

```
brew install kubectl
brew install minikube
```

Verify installations:

```
kubectl --help
minikube --help
```

Test on Local Machine

1. Set Up the Environment

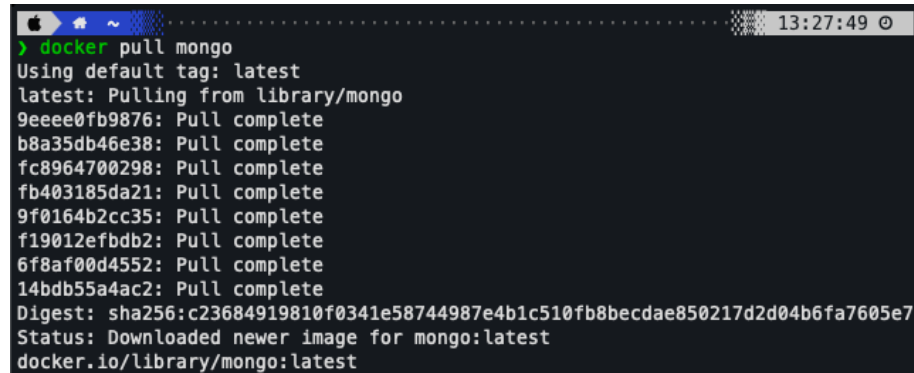
Unzip the downloaded file, navigate into the project folder, and install dependencies:

```
pip install -r requirements.txt
```

2. Run MongoDB Using Docker

Pull the official MongoDB image:

```
docker pull mongo
```



```
> docker pull mongo
Using default tag: latest
latest: Pulling from library/mongo
9eeee0fb9876: Pull complete
b8a35db46e38: Pull complete
fc8964700298: Pull complete
fb403185da21: Pull complete
9f0164b2cc35: Pull complete
f19012efbdb2: Pull complete
6f8af00d4552: Pull complete
14bdb55a4ac2: Pull complete
Digest: sha256:c23684919810f0341e58744987e4b1c510fb8becdae850217d2d04b6fa7605e7
Status: Downloaded newer image for mongo:latest
docker.io/library/mongo:latest
```

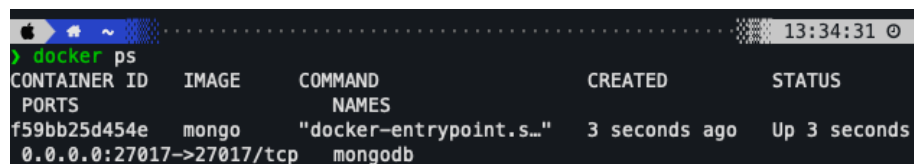
Run the MongoDB container:

```
docker run -d --name mongodb -p 27017:27017 -v mongo_data:/data/db mongo
```

Explanation: - d → Run in detached mode (background) - --name mongodb → Container name - -p 27017:27017 → Expose MongoDB default port - -v mongo_data:/data/db → Persist data using Docker volume - mongo → Use the official MongoDB image

Check if the container is running:

```
docker ps
```



```
> docker ps
CONTAINER ID   IMAGE     COMMAND                  CREATED        STATUS
PORTS
f59bb25d454e   mongo    "docker-entrypoint.s..." 3 seconds ago  Up 3 seconds
0.0.0.0:27017->27017/tcp    mongodb
```

Figure 1: Docker PS Mongo

3. MongoDB Configuration in app.py

The Flask app connects to MongoDB as follows:

```
mongodb_host = os.environ.get('MONGO_HOST', 'localhost')
mongodb_port = int(os.environ.get('MONGO_PORT', '27017'))
client = MongoClient(mongodb_host, mongodb_port)
db = client.camp2016
todos = db.todo
```

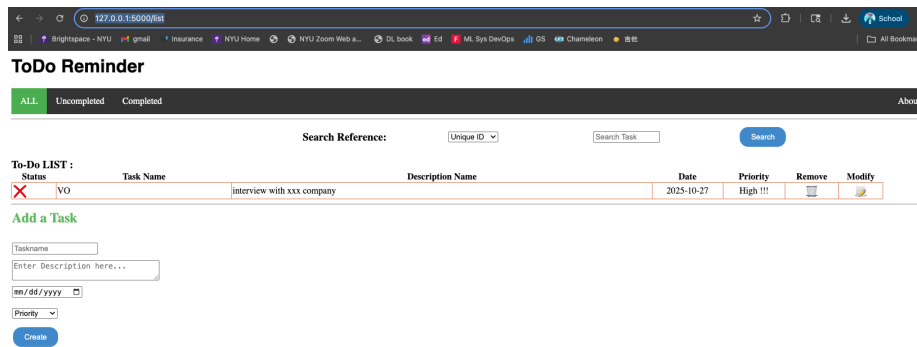
4. Start the Flask Application

Run the Flask app:

```
python3 app.py
```

Then open your browser and go to:

`http://127.0.0.1:5000`



to verify the application is running locally.

Inspect Docker Volumes

To check your MongoDB volume:

```
docker volume inspect mongo_data
docker volume ls
```

```

> docker volume inspect mongo_data
[
  {
    "CreatedAt": "2025-10-14T17:34:31Z",
    "Driver": "local",
    "Labels": null,
    "Mountpoint": "/var/lib/docker/volumes/mongo_data/_data",
    "Name": "mongo_data",
    "Options": null,
    "Scope": "local"
  }
]

> docker volume ls
DRIVER      VOLUME NAME
local       ad5e9c4af71df7749ded8a0adb8ac6d819e951c9bccda34615133fcb9fb3c748
local       mongo_data

```

Inspect MongoDB Data

Access MongoDB shell inside the container:

```
docker exec -it mongodb mongosh
```

Then run:

```
use camp2016           # switch to database
show collections        # list collections
db.todo.find().pretty() # view data inside 'todo' collection
```

You've successfully set up Docker, MongoDB, and Flask locally!

Containize

1.Create a Dockfile for Flask App

```
# Use the official Python 3.8 slim image as the base image
FROM python:3.8-slim

# Set the working directory within the container
WORKDIR /app

COPY requirements.txt .

# Upgrade pip and install Python dependencies
RUN pip3 install --upgrade pip && pip install --no-cache-dir -r requirements.txt

# Copy the rest of app
COPY . .
```

```
# Expose Flask port default
EXPOSE 5000
```

```
CMD ["python", "app.py"]
```

2. build the image

```
docker build -t flask-app .
```

3. build the container

```
docker run -d -p 5000:5000 flask-app flask-app
```

bug

- MacOS using port 5000 for airplay, need to shut down it to use 5000 port
 - pymongo.errors.ServerSelectionTimeoutError: localhost:27017: [Errno 111] Connection refused
 - inside the flask container, localhost means the Flask container itself, not the host or Mongo container.
- ```
client = MongoClient('localhost', 27017)
```

### fix it

- using docker-compose.yml
- for localtest manual docker run with --network

#create a bridge network name flasknet, so containers on the same network can communicate by

```
docker network create flasknet
```

```
docker run -d --name mongodb --network flasknet -p 27017:27017 mongo
```

```
docker run -d --name flask-app --network flasknet -p 5000:5000 flasknet -e MONGO_HOST=mongo
```

## 4. push to Docker Hub

we can create a new tag and push the image to the Docker Hub

```
docker push majiny/flask-app:latest
```

## Deploy to the minikube

Since we have minikube downloaded before, so that we can create the deploy to the minikube.

### 1. Start minikube

using the doc to start, minikube service

```
minikube start
```

```
~/De/N/NYU_Fall_2025/CS-GY_9/pr/Cloud - Assignment 3 Files [main] 12:02:11
$ docker tag flask-app:latest majiny/flask-app:latest

~/De/N/NYU_Fall_2025/CS-GY_9/pr/Cloud - Assignment 3 Files [main] 12:02:22
$ ls
Dockerfile app.py requirements.txt static venv
README.md docker script templates

~/De/N/NYU_Fall_2025/CS-GY_9/pr/Cloud - Assignment 3 Files [main] 12:02:23
$ docker images
REPOSITORY TAG IMAGE ID CREATED SIZE
flask-app latest e122ce09f3e8 3 minutes ago 282MB
majiny/flask-app latest e122ce09f3e8 3 minutes ago 282MB
mongo latest c23664919810 3 weeks ago 1.2GB

~/De/N/NYU_Fall_2025/CS-GY_9/pr/Cloud - Assignment 3 Files [main] 12:02:30
$ docker push majiny/flask-app:latest
The push refers to repository [docker.io/majiny/flask-app]
442bf64da915: Pushed
14c9d9d19932: Pushed
d847ad1879f2: Pushed
70414efc397c: Pushed
9b5da5950b8f: Pushed
d8f14ea64647: Pushed
56d45b8e1404: Pushed
6cd45dd32d5: Pushed
b77c01107b73: Pushed
latest: digest: sha256:e122ce09f3e8d6d1c2f8fed013592af23fd85bf26705d0c01eacb820685ab5b5 size: 856
```

Figure 2: Push to DockerHub

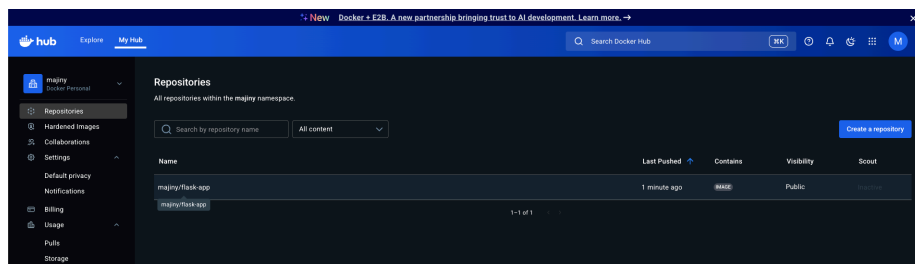


Figure 3: Image in DockerHub

## 2. create the yaml file

### 1. mongodb-pvc.yaml - Persistent Storage for MongoDB:

- persistentVolumeClaim(PVC)-> A request for storage space in Kubernetes.
- accessModes: ReadWriteOnce-> The volumen can be mounted as read-write by only one node.
- resources.requests.storage-> Requests 256 MB of space.,

MongoDB stores its database files under /data/db. Without a PVC, that data disappears when t

### 2. mongodb-deployment.yaml - Running MongoDB Pod:

- kind: Deployment-> Ensures a specific # of identical pods run continuously.
- metadata.name-> Deployment name (mongodb-deployment).
- selector/labels-> Used to connect pods, services, and deployments.
- containers-> Defines which image to run (we used mongo:latest).
- containerPort-> MongoDB listens on port 27017 (default port for MongoDB).
- aresources.requests-> Reserves 256 MB memory and 0.5 CPU for Mongo.
- volumeMounts-> Mounts the PVC inside the container at /data/db.
- volumes-> Attaches the PVC created earlier (mongodb-pvc)

MongoDB Deployment spins up MongoDB and mounts a persistent volume for storage. Even if the

### 3. mongodb-service.yaml - Internal Service for MongoDB:

- kind: Service-> exposes a stable endpoint (DNS name) for Mongo pods.
- selector.app: mongo-> Connects this Service to pods with label app: mongo.
- port/ targetPort-> Routes traffic from Service port 27017-> pod port 27017.

Flask can not directly know a pod's IP (pods are dynamic). Instead, it connects to mongo-ser

### 4. flask-deployment.yaml - Flask App Deployment:

- replicas: 1 -> Run on Flask pod (for Minikube demon).
- images: majiny/flask-app-> The flask app we pushed to Docker Hub before.
- containerPort: 5000 -> Flask runs internally on port 5000.
- en variables -> Tell Flask where to find MongoDB.
  - MONGO\_HOST=mongo-service
  - MONGO\_PORT=27017
- imagePullPolicy: IfNotPresent->Uses a cached local image if available

This defines the applicaiton container. It runs Flask, connects to MongoDB vai service name.

### 5. flask-service.yaml - Exposing Flask to the Outside World:

- type: LoadBalancer -> Create an external access point/ outside cluster.
- port: 5000 -> Port exposed to the world.
- targetPort: 5000 -> Port inside the pod where Flaks runs.
- selector.app: to-do \_> Link to Flaks pod labeld app: to-do

This lets users access ahte Flask app from outside Kubernetes.

## 3. deploy service

using the code to deploy the service.

```
kubect1 apply -f mongodb-pvc.yaml
kubect1 apply -f mongodb-service.yaml
kubect1 apply -f mongodb-deployment.yaml
kubect1 apply -f flask-deployment.yaml
kubect1 apply -f flask-service.yaml
```

```

$ kubectl apply -f mongodb-pvc.yaml
$ kubectl apply -f mongodb-service.yaml
$ kubectl apply -f mongodb-deployment.yaml
$ kubectl apply -f flask-deployment.yaml
$ kubectl apply -f flask-service.yaml
persistentvolumeclaim/mongo-pvc created
service/mongo-service created
deployment.apps/mongo-deployment created
deployment.apps/to-do-deployment created
service/to-do-service created

$ kubectl get all
NAME READY STATUS RESTARTS AGE
pod/mongo-deployment-67cc7ddf6b-7qqln 1/1 Running 0 19s
pod/to-do-deployment-79db69858d-89jmv 1/1 Running 0 19s

NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
service/kubernetes ClusterIP 10.96.0.1 <none> 443/TCP 26h
service/mongo-service ClusterIP 10.106.29.130 <none> 27017/TCP 19s
service/to-do-service LoadBalancer 10.107.233.170 <pending> 5000:31154/TCP 19s

NAME READY UP-TO-DATE AVAILABLE AGE
deployment.apps/mongo-deployment 1/1 1 1 19s
deployment.apps/to-do-deployment 1/1 1 1 19s

NAME DESIRED CURRENT READY AGE
replicaset.apps/mongo-deployment-67cc7ddf6b 1 1 1 19s
replicaset.apps/to-do-deployment-79db69858d 1 1 1 19s

```

Figure 4: Deploy On Minikube

#### 4. Test the app in Minikube

To ensure it is running on Minikube by accessing it using the service URL  
 minikube service to-do-service

```

$ minikube service to-do-service

```

| NAMESPACE | NAME          | TARGET PORT | URL                       |
|-----------|---------------|-------------|---------------------------|
| default   | to-do-service | 5000        | http://192.168.49.2:31154 |

```

Starting tunnel for service to-do-service.

```

| NAMESPACE | NAME          | TARGET PORT | URL                    |
|-----------|---------------|-------------|------------------------|
| default   | to-do-service |             | http://127.0.0.1:64879 |

```

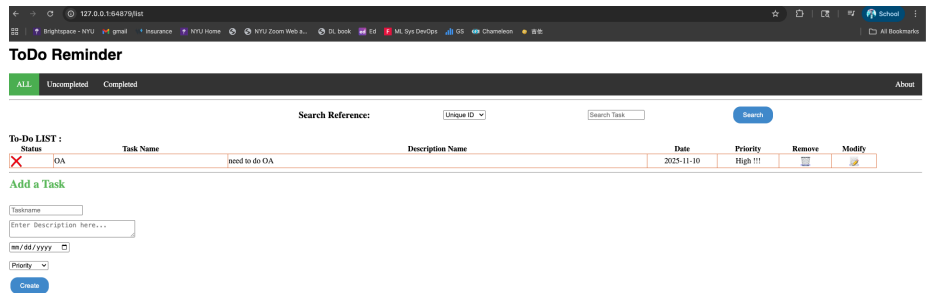
Starting tunnel for service to-do-service.
Opening service default/to-do-service in default browser...
Because you are using a Docker driver on darwin, the terminal needs to be open to run it.

```

Figure 5: Minikube service

and we can see the URL address is <http://127.0.0.1:64879> and we can visited the





site by using this URL. .

## Deploy to AWS EKS

### 1. download AWS CLI

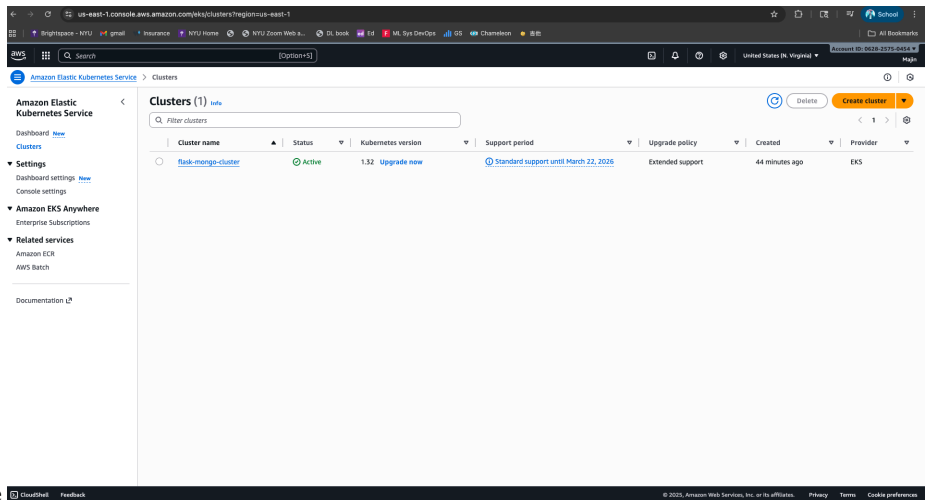
```
curl "https://awscli.amazonaws.com/AWSCLIV2.pkg" -o "AWSCLIV2.pkg"
sudo installer -pkg AWSCLIV2.pkg -target /
```

```
$ which aws
/usr/local/bin/aws
$ aws --version
aws-cli/2.31.9 Python/3.13.7 Darwin/25.1.0 source/arm64
```

### 2. Create cluster

the easy we to create the cluster, we can use [eksctl][<https://docs.aws.amazon.com/eks/latest/userguide/getting-started-eksctl.html>]

```
eksctl create cluster \
 --name flask-mongo-cluster \
 --region us-east-1 \
 --nodes 1 \
 --node-type t3.micro
```



After launch, we can see

since there we limited cpu as 500m and memory: 256Mi. and when we trying to deploy the app, there in no enough resource, so it in the pending state

```

$ kubectl get all
NAME READY STATUS RESTARTS AGE
pod/mongo-deployment-6d5f989dcc-qws6 0/1 Pending 0 113s
pod/to-do-deployment-7cb67d485c-ktxf 0/1 Pending 0 112s

NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
service/kubernetes ClusterIP 10.100.0.1 <none> 443/TCP 14m
service/mongo-service ClusterIP 10.100.186.163 <none> 27017/TCP 113s
service/to-do-service LoadBalancer 10.100.171.36 a479634b047bf4ff89de68da6e02f69a-1169876625.us-east-1.elb.amazonaws.com 5000:30690/TCP 112s

NAME READY UP-TO-DATE AVAILABLE AGE
deployment.apps/mongo-deployment 0/1 1 0 113s
deployment.apps/to-do-deployment 0/1 1 0 112s

NAME DESIRED CURRENT READY AGE
replicaset.apps/mongo-deployment-6d5f989dcc 1 1 0 113s
replicaset.apps/to-do-deployment-7cb67d485c 1 1 0 112s

```

using the kubectl to check why it is pending

`kubectl describe pod <POD_ID>`

delete the resource we had before and rebuild

`eksctl delete cluster --name flask-mongo-cluster --region us-east-1`

```

eksctl create cluster \
 --name flask-mongo-cluster \
 --region us-east-1 \
 --nodegroup-name flask-nodegroup \
 --nodes 2 \
 --nodes-min 1 \
 --nodes-max 3 \
 --managed \
 --node-type t3.small

```

```

$ kubectl describe pod mongo-deployment-6d5f989dcd-wqws6
Name: mongo-deployment-6d5f989dcd-wqws6
Namespace: default
Priority: 0
Service Account: default
Node: ip-10-0-1-10.us-east-1.compute.amazonaws.com
Labels: app=mongo, pod-template-hash=6d5f989dcd
Annotations: <none>
Status: Pending
IP: <none>
IPNotReady: <none>
Controlled By: ReplicaSet/mongo-deployment-6d5f989dcd
Containers:
 mongo:
 Image: mongo:latest
 Port: 27017/TCP
 Host Port: 27017/TCP
 Requests: cpu=500m, memory=250Mi
 Environment: <none>
 Mounts:
 /data/db from storage (rw)
 /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-58vsr (ro)
Conditions:
 Type Status
 PodScheduled False
Volumes:
 storage:
 Type: PersistentVolumeClaim (a reference to a PersistentVolumeClaim in the same namespace)
 ClaimName: mongo-pvc
 ReadOnly: false
 kube-api-access-58vsr:
 Type: Projected (a volume that contains injected data from multiple sources)
 TokenExpirationSeconds: 3607
 ConfigMapName: kube-root-ca.crt
 Optional: false
 DownwardAPI: true
QoS Class: Burstable
Node-Selectors: <none>
Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
 node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
 Type Reason Age From Message
 ---- ------ -
Warning FailedScheduling 2m28s default-scheduler 0/1 nodes are available: 1 Insufficient memory, 1 Too many pods, preemption: 0/1 nodes are available: 1 No preemption victims found for incoming pod.

```

Figure 6: kubectl\_describe

### 3. set up addon

Since we want to use pvc storage we need to Enable OIDC provbider (required for IAM-> K8s mapping)

```
eksctl utils associate-iam-oidc-provider \
 --region us-east-1 \
 --cluster flask-mongo-cluster \
 --approve
```

then create IAM role for the EBS CSI driver

```
eksctl create iamserviceaccount \
 --region us-east-1 \
 --name ebs-csi-controller-sa \
 --namespace kube-system \
 --cluster flask-mongo-cluster \
 --attach-policy-arn arn:aws:iam::aws:policy/service-role/AmazonEBSCSIDriverPolicy \
 --approve \
 --role-name AmazonEKS_EBS_CSI_DriverRole
```

Install the AWS EBS CSI driver add-on

```
aws eks create-addon \
 --region us-east-1 \
 --cluster-name flask-mongo-cluster \
 --addon-name aws-ebs-csi-driver \
 --service-account-role-arn arn:aws:iam::<id>:role/AmazonEKS_EBS_CSI_DriverRole #use the
```

To verify installation

```
kubectl get pods -n kube-system | grep ebs
```

```
> kubectl get pods -n kube-system | grep ebs
ebs-csi-controller-5d97d64f67-gqkfh 6/6 Running 0 3m43s
ebs-csi-controller-5d97d64f67-qvltw 6/6 Running 0 3m43s
ebs-csi-node-2k2dt 3/3 Running 0 3m43s
ebs-csi-node-bhpm4 3/3 Running 0 3m43s
```

### 4. deploy to eks To get pvc running we run same code we had as we deploy to the Minikube

```
kubectl apply -f mongodb-pvc.yaml
kubectl get pvc
```

```
~/De/N/NYU_Fall_2025/CS-GY 9/pr/Cloud - Assignment 3 Files/docker main !3
> kubectl apply -f mongodb-pvc.yaml
persistentvolumeclaim/mongo-pvc created

~/De/N/NYU_Fall_2025/CS-GY 9/pr/Cloud - Assignment 3 Files/docker main !3
> kubectl get pvc
error: the server doesn't have a resource type "pvc"

~/De/N/NYU_Fall_2025/CS-GY 9/pr/Cloud - Assignment 3 Files/docker main !3
> kubectl get pvc
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS VOLUMEATTRIBUTESCLASS AGE
mongo-pvc Pending gp2 <unset> 9s
```

we will see there is still no VOLUME, CAPACITY, ACCESS, and MODES by checking pvc mongo-pvc

```
kubectl describe pvc mongo-pvc
```

we see the message waiting for first Consumer to be created before binding

```
~/De/N/NYU_Fall_2025/CS-GY 9/pr/Cloud - Assignment 3 Files/docker main !3
> kubectl describe pvc mongo-pvc
Name: mongo-pvc
Namespace: default
StorageClass: gp2
Status: Pending
Volume:
Labels: <none>
Annotations: <none>
Finalizers: [kubernetes.io/pvc-protection]
Capacity:
Access Modes:
VolumeMode: Filesystem
Used By: <none>
Events:
Type Reason Age From Message
Normal WaitForFirstConsumer 14s (x6 over 80s) persistentvolume-controller waiting for first consumer to be created before binding
```

After we apply mongo-deployment

```
kubectl apply -f mongodbv-deployment.yaml
kubectl get pods
kubectl get pvc
```

```
~/De/N/NYU_Fall_2025/CS-GY 9/pr/Cloud - Assignment 3 Files/docker main !3
> kubectl apply -f mongodbv-deployment.yaml
deployment.apps/mongo-deployment created

~/De/N/NYU_Fall_2025/CS-GY 9/pr/Cloud - Assignment 3 Files/docker main !3
> kubectl get pods
NAME READY STATUS RESTARTS AGE
mongo-deployment-6d5f989dcc-jqzb5 0/1 ContainerCreating 0 7s

~/De/N/NYU_Fall_2025/CS-GY 9/pr/Cloud - Assignment 3 Files/docker main !3
> kubectl get pvc
NAME STATUS VOLUME CAPACITY ACCESS MODES STORAGECLASS VOLUMEATTRIBUTESCLASS AGE
mongo-pvc Bound pvc-c3b6222-f140-4683-a793-511a7aacd207 1Gi RWO gp2 <unset> 118s
```

Figure 7: complete PVC

then we apply reset of yaml file

```
kubectl apply -f mongodb-service.yaml
kubectl apply -f flask-deployment.yaml
kubectl apply -f flask-service.yaml
kubectl get all
```

```
kubectl get all
```

| NAME                                  | READY | STATUS  | RESTARTS | AGE |
|---------------------------------------|-------|---------|----------|-----|
| pod/mongo-deployment-6d5f989dcc-jqzb5 | 1/1   | Running | 0        | 91s |
| pod/to-do-deployment-7cb67d485c-5vnnk | 1/1   | Running | 0        | 7s  |

| NAME                  | TYPE         | CLUSTER-IP     | EXTERNAL-IP                                                             | PORT(S)        | AGE |
|-----------------------|--------------|----------------|-------------------------------------------------------------------------|----------------|-----|
| service/kubernetes    | ClusterIP    | 10.100.0.1     | <none>                                                                  | 443/TCP        | 32m |
| service/mongo-service | ClusterIP    | 10.100.38.61   | <none>                                                                  | 27017/TCP      | 65s |
| service/to-do-service | LoadBalancer | 10.100.191.175 | a0a352132be6044f88cc012031f6faa3-1095562976.us-east-1.elb.amazonaws.com | 5000:30101/TCP | 15s |

| NAME                             | READY | UP-TO-DATE | AVAILABLE | AGE |
|----------------------------------|-------|------------|-----------|-----|
| deployment.apps/mongo-deployment | 1/1   | 1          | 1         | 91s |
| deployment.apps/to-do-deployment | 1/1   | 1          | 1         | 7s  |

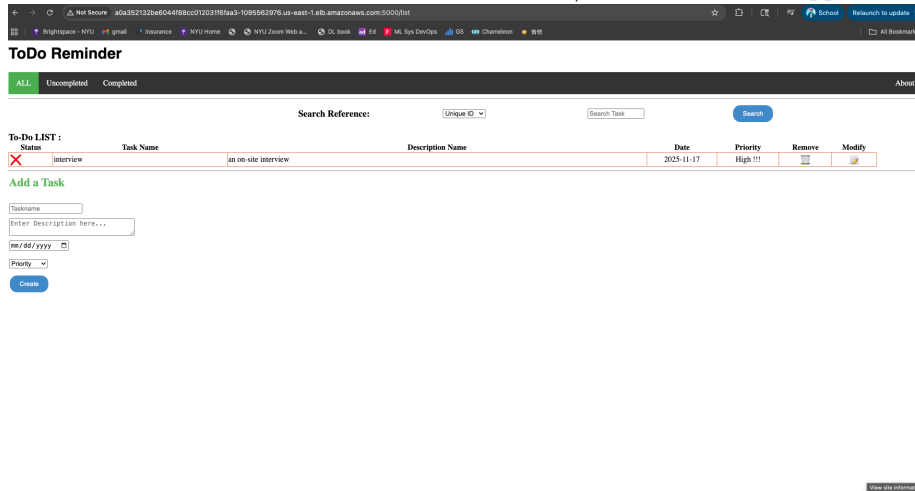
  

| NAME                                        | DESIRED | CURRENT | READY | AGE |
|---------------------------------------------|---------|---------|-------|-----|
| replicaset.apps/mongo-deployment-6d5f989dcc | 1       | 1       | 1     | 91s |
| replicaset.apps/to-do-deployment-7cb67d485c | 1       | 1       | 1     | 7s  |

Figure 8: k8s\_aws

The all nodes are running, and we have a external-ip for our to-do-service but we can not visited yet, it takes time to set up waiting\_external\_ip

now we can use <http://a0a352132be6044f88cc012031f6faa3-1095562976.us-east-1.elb.amazonaws.com:5000/> to visited our application



## Deployments and ReplicaSets

Before we change the replicas1 we have defined replicas as 1

```
spec:
 replicas: 1
```

after we change it to 5

```
spec:
```

```

$ kubectl get all
NAME READY STATUS RESTARTS AGE
pod/mongo-deployment-6d5f989dcc-jqzb5 1/1 Running 0 31m
pod/to-do-deployment-7cb67d485c-5vnnk 1/1 Running 0 30m

NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
service/kubernetes ClusterIP 10.100.0.1 <none> 443/TCP 63m
service/mongo-service ClusterIP 10.100.38.61 <none> 27017/TCP 31m
service/to-do-service LoadBalancer 10.100.191.175 a0a352132be6044f88cc012031f6faa3-1095562976.us-east-1.elb.amazonaws.com 5000:30101/TCP 30m

NAME READY UP-TO-DATE AVAILABLE AGE
deployment.apps/mongo-deployment 1/1 1 1 31m
deployment.apps/to-do-deployment 1/1 1 1 30m

NAME DESIRED CURRENT READY AGE
replicaset.apps/mongo-deployment-6d5f989dcc 1 1 1 31m
replicaset.apps/to-do-deployment-7cb67d485c 1 1 1 30m

```

Figure 9: get\_all\_re1

replicas: 5

```

kubectl apply -f flask-deployment.yaml
Verify the ReplicasSet
kubectl get all
kubectl get rs

```

```

$ kubectl get all
NAME READY STATUS RESTARTS AGE
pod/mongo-deployment-6d5f989dcc-jqzb5 1/1 Running 0 42m
pod/to-do-deployment-7cb67d485c-5vnnk 1/1 Running 0 41m
pod/to-do-deployment-7cb67d485c-6c29d 1/1 Running 0 10s
pod/to-do-deployment-7cb67d485c-jh2dg 1/1 Running 0 10s
pod/to-do-deployment-7cb67d485c-lmddp 1/1 Running 0 10s
pod/to-do-deployment-7cb67d485c-p5sbw 1/1 Running 0 10s

NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
service/kubernetes ClusterIP 10.100.0.1 <none> 443/TCP 73m
service/mongo-service ClusterIP 10.100.38.61 <none> 27017/TCP 42m
service/to-do-service LoadBalancer 10.100.191.175 a0a352132be6044f88cc012031f6faa3-1095562976.us-east-1.elb.amazonaws.com 5000:30101/TCP 41m

NAME READY UP-TO-DATE AVAILABLE AGE
deployment.apps/mongo-deployment 1/1 1 1 42m
deployment.apps/to-do-deployment 5/5 5 5 41m

NAME DESIRED CURRENT READY AGE
replicaset.apps/mongo-deployment-6d5f989dcc 1 1 1 42m
replicaset.apps/to-do-deployment-7cb67d485c 5 5 5 41m

$ kubectl get rs
NAME DESIRED CURRENT READY AGE
mongo-deployment-6d5f989dcc 1 1 1 42m
to-do-deployment-7cb67d485c 5 5 5 41m

```

Figure 10: Verify ReplicasSet

When we delete the pod we need to know the pod name

```

kubectl get pods
kubectl delete pod to-do-deployment-7cb67d485c-5vnnk to-do-deployment-7cb67d485c-6c29d
kubectl get pods

```

we will see the two pods were deleted, and by checking pods we will see two new

```
~/De/N/NYU_Fall_2025/CS-GY 9/pr/Cloud - Assignment 3 Files/docker P main !4 ?12
> kubectl get pods
NAME READY STATUS RESTARTS AGE
mongo-deployment-6d5f989dcc-jqzb5 1/1 Running 0 47m
to-do-deployment-7cb67d485c-5vnnk 1/1 Running 0 46m
to-do-deployment-7cb67d485c-6c29d 1/1 Running 0 5m21s
to-do-deployment-7cb67d485c-jh2dg 1/1 Running 0 5m21s
to-do-deployment-7cb67d485c-lmddp 1/1 Running 0 5m21s
to-do-deployment-7cb67d485c-p5sbw 1/1 Running 0 5m21s

~/De/N/NYU_Fall_2025/CS-GY 9/pr/Cloud - Assignment 3 Files/docker P main !4 ?13
> kubectl delete pod to-do-deployment-7cb67d485c-5vnnk to-do-deployment-7cb67d485c-6c29d
pod "to-do-deployment-7cb67d485c-5vnnk" deleted from default namespace
pod "to-do-deployment-7cb67d485c-6c29d" deleted from default namespace

~/De/N/NYU_Fall_2025/CS-GY 9/pr/Cloud - Assignment 3 Files/docker P main !4 ?13
> kubectl get pods
NAME READY STATUS RESTARTS AGE
mongo-deployment-6d5f989dcc-jqzb5 1/1 Running 0 48m
to-do-deployment-7cb67d485c-9gtjd 1/1 Running 0 9s
to-do-deployment-7cb67d485c-fddc5 1/1 Running 0 9s
to-do-deployment-7cb67d485c-jh2dg 1/1 Running 0 6m31s
to-do-deployment-7cb67d485c-lmddp 1/1 Running 0 6m31s
to-do-deployment-7cb67d485c-p5sbw 1/1 Running 0 6m31s
```

pods created

we can scale down the replicas without changing yaml file

`kubectl scale deployment to-do-deployment --replicas=3`

```
~/De/N/NYU_Fall_2025/CS-GY 9/pr/Cloud - Assignment 3 Files/docker P main !4 ?14
> kubectl scale deployment --replicas=3
error: resource(s) were provided, but no name was specified

~/De/N/NYU_Fall_2025/CS-GY 9/pr/Cloud - Assignment 3 Files/docker P main !4 ?14
> kubectl scale deployment to-do-deployment --replicas=3
deployment.apps/to-do-deployment scaled

~/De/N/NYU_Fall_2025/CS-GY 9/pr/Cloud - Assignment 3 Files/docker P main !4 ?14
> kubectl get pods
NAME READY STATUS RESTARTS AGE
mongo-deployment-6d5f989dcc-jqzb5 1/1 Running 0 53m
to-do-deployment-7cb67d485c-jh2dg 1/1 Running 0 11m
to-do-deployment-7cb67d485c-lmddp 1/1 Running 0 11m
to-do-deployment-7cb67d485c-p5sbw 1/1 Running 0 11m
```

Figure 11: CLI scaled down